

# Computer Memory

Dr. S Pal, CSE, IIT(BHU)


Slide Sources: Patterson &

Hennessy COD book site

(copyright Morgan Kaufmann)

adapted and supplemented

# Five basic components of a computer




# Memories: Review

## *SRAM (Static Random Access Memory):*

- value is stored on a pair of inverting gates that will *exist indefinitely* as long as there is power, which is why it is called *static*
- very fast but takes up more space than DRAM – 4 to 6 transistors per bit
- used for *cache*

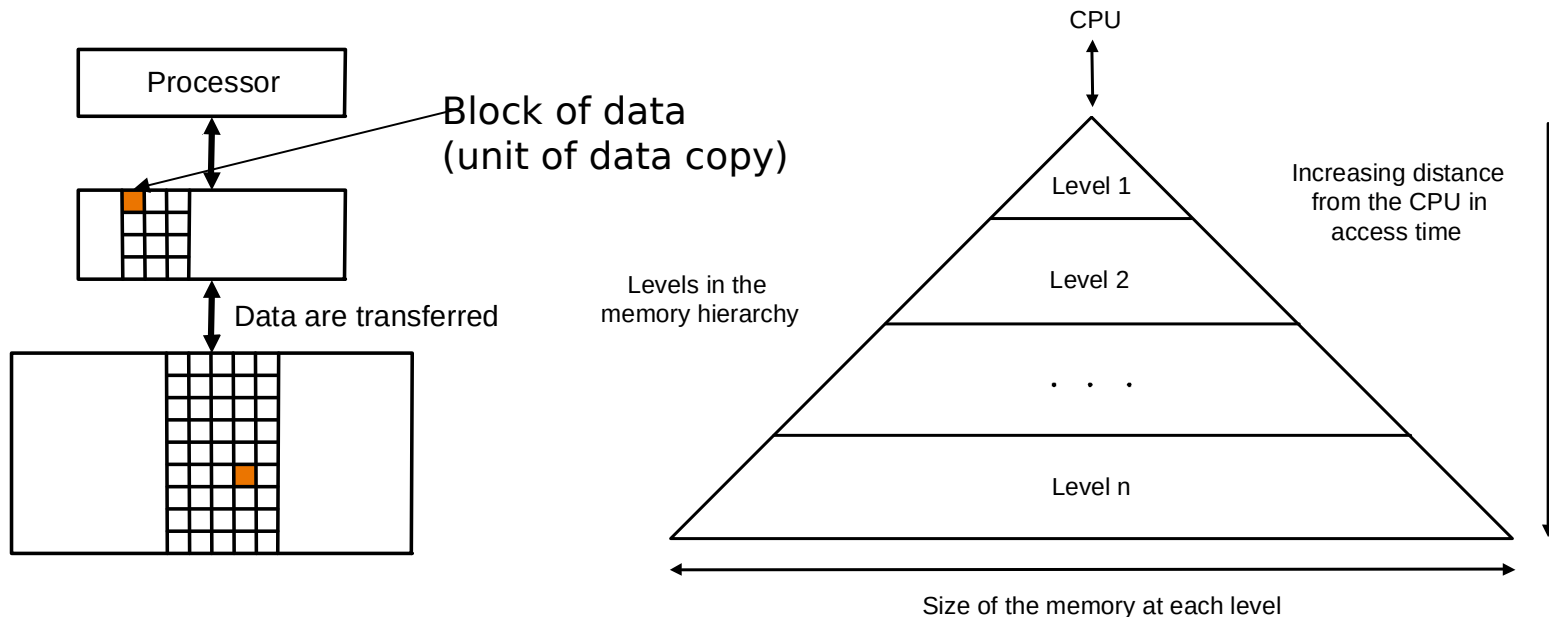
## *DRAM (Dynamic Random Access Memory):*

- value is stored as a charge on capacitor that must be *periodically refreshed*, which is why it is called *dynamic*
- very small – 1 transistor per bit – but factor of 5 to 10 slower than SRAM
- used for *main memory*



# Memory Hierarchy

- Users want **large** and **fast memories**...
  - expensive and they don't like to pay...
- Make it seem like they have what they want...
  - *memory hierarchy*
  - hierarchy is *inclusive*, every level is *subset* of lower level
  - performance depends on *hit rates*





# Locality

---

- *Locality* is a principle that makes having a memory hierarchy a good idea
- If an item is referenced then because of
  - *temporal locality*: it will tend to be *again* referenced soon
  - *spatial locality*: *nearby items* will tend to be referenced soon
  - *why does code have locality – consider instruction and data?*



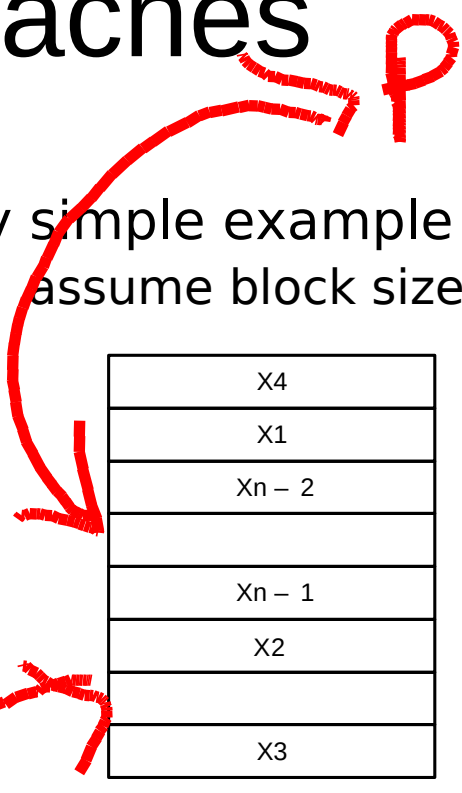
# Hit and Miss

---

- Focus on *any two adjacent* levels – called, *upper* (closer to CPU) and *lower* (farther from CPU) – in the memory hierarchy, because each block copy is always between two adjacent levels
- Terminology:
  - *block*: minimum unit of data to move between levels
  - *hit*: data requested is available in upper level
  - *miss*: data requested is not available in upper level
  - *hit rate*: fraction of memory accesses that are hits (i.e., found at upper level)
  - *miss rate*: fraction of memory accesses that are not hits
    - $\text{miss rate} = 1 - \text{hit rate}$
  - *hit time*: time to determine if the access is indeed a hit + time to access and deliver the data from the upper level to the CPU
  - *miss penalty*: time to determine if the access is a miss + time to replace block at upper level with corresponding block at lower level + time to deliver the block to the CPU

# Caches

- By simple example
  - assume block size = one word of data



X4
X1
Xn - 2
Xn - 1
X2
X3

a. Before the reference to X<sub>n</sub>

X4
X1
Xn - 2
Xn - 1
X2
X <sub>n</sub>
X3

b. After the reference to X<sub>n</sub>


Reference to X<sub>n</sub>  
causes miss so  
it is fetched from  
memory

- Issues:
  - *how do we know if a data item is in the cache?*
  - *if it is, how do we find it?*
  - *if not, what do we do?*

- Solution depends on *cache addressing scheme...*

# Direct Mapped Cache

- Addressing scheme in *direct mapped* cache:
  - cache block address = memory block address *mod* cache size (*unique*)
  - if cache size =  $2^m$ , cache address = lower **m** bits of **n**-bit memory address
  - remaining upper (**n-m**) bits kept as *tag bits* at each cache block
  - also need a *valid bit* to recognize valid entry





# Accessing Cache

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	$10110_{\text{two}}$	miss (5.9b)	$(10110_{\text{two}} \bmod 8) = 110_{\text{two}}$
26	$11010_{\text{two}}$	miss (5.9c)	$(11010_{\text{two}} \bmod 8) = 010_{\text{two}}$
22	$10110_{\text{two}}$	hit	$(10110_{\text{two}} \bmod 8) = 110_{\text{two}}$
26	$11010_{\text{two}}$	hit	$(11010_{\text{two}} \bmod 8) = 010_{\text{two}}$
16	$10000_{\text{two}}$	miss (5.9d)	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$
3	$00011_{\text{two}}$	miss (5.9e)	$(00011_{\text{two}} \bmod 8) = 011_{\text{two}}$
16	$10000_{\text{two}}$	hit	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$
18	$10010_{\text{two}}$	miss (5.9f)	$(10010_{\text{two}} \bmod 8) = 010_{\text{two}}$
16	$10000_{\text{two}}$	hit	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$

(0) Initial state:

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

a. The initial state of the cache after power-on

(1) Address referred  $10110$  (*miss*):

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	$10_{\text{two}}$	Memory ( $10110_{\text{two}}$ )
111	N		

b. After handling a miss of address ( $10110_{\text{two}}$ )

Index	V	Tag	Data
000	N		
001	N		
010	Y	$11_{\text{two}}$	Memory ( $11010_{\text{two}}$ )
011	N		
100	N		
101	N		
110	Y	$10_{\text{two}}$	Memory ( $10110_{\text{two}}$ )
111	N		

c. After handling a miss of address ( $11010_{\text{two}}$ )

Index	V	Tag	Data
000	Y	$10_{\text{two}}$	Memory ( $10000_{\text{two}}$ )
001	N		
010	Y	$11_{\text{two}}$	Memory ( $11010_{\text{two}}$ )
011	Y	$00_{\text{two}}$	Memory ( $00011_{\text{two}}$ )
100	N		
101	N		
110	Y	$10_{\text{two}}$	Memory ( $10110_{\text{two}}$ )
111	N		

e. After handling a miss of address ( $00011_{\text{two}}$ )


Index	V	Tag	Data
000	Y	$10_{\text{two}}$	Memory ( $10000_{\text{two}}$ )
001	N		
010	Y	$11_{\text{two}}$	Memory ( $11010_{\text{two}}$ )
011	N		
100	N		
101	N		
110	Y	$10_{\text{two}}$	Memory ( $10110_{\text{two}}$ )
111	N		

d. After handling a miss of address ( $10000_{\text{two}}$ )

Index	V	Tag	Data
000	Y	$10_{\text{two}}$	Memory ( $10000_{\text{two}}$ )
001	N		
010	Y	$10_{\text{two}}$	Memory ( $10010_{\text{two}}$ )
011	Y	$00_{\text{two}}$	Memory ( $00011_{\text{two}}$ )
100	N		
101	N		
110	Y	$10_{\text{two}}$	Memory ( $10110_{\text{two}}$ )
111	N		

f. After handling a miss of address ( $10010_{\text{two}}$ )

# Direct Mapped Cache



**Cache with 1024 1-word blocks: *byte offset* (least 2 significant bits) is ignored and next 10 bits used to index into cache**

# Cache arithmetic

- 64-bit addresses
- A direct-mapped cache
- The cache size is  $2^n$  blocks, so  $n$  bits are used for the index
- The block size is  $2^m$  words ( $2^{m+2}$  bytes), so  $m$  bits are used for the word within the block, and 2 bits are used for the byte part of the address
- The size of the tag field is =  $64 - (n+m+2)$
- Total #bits  $2^n \times (\text{block size} + \text{tag size} + \text{valid field size})$   
$$= 2^n \times (2^m \times 32 + (64 - n - m - 2) + 1) = 2^n \times (2^m \times 32 + 63 - n - m).$$
- But *tag*, *valid* bits are ignored when cache size is referred (here, 4KB cache)

**Example:** How many total bits are required for a direct-mapped cache with 16 KiB of data and ~~1-word~~ blocks, assuming a 64-bit address?

**Ans:** 16 KiB is  $2^4 \times 2^{10}$  bytes =  $2^{12}$  words

With a block size of 1 word, there are 4096 ( $2^{12}$ ) blocks

Each block has 1 word = 32 bits of data plus a tag, which is  $(64 - 12 - 2)$  bits, plus a valid bit.

Thus, the complete cache size is

$$= 2^{12} \times \{32 + (64 - 12 - 2) + 1\} = 2^{12} \times 83 \text{ bits}$$

= **33.2 KB** for a ~~16 KB~~ cache.

**Example:** How many total bits are required for a direct-mapped cache with 16 KiB of data and four-word blocks, assuming a 64-bit address?

**Ans:** 16 KiB is 4096 ( $2^{12}$ ) words

With a block size of 4 words ( $2^2$ ), there are 1024 ( $2^{10}$ ) blocks

Each block has  $4 \times 32 = 128$  bits of data plus a tag, which is  $(64 - 10 - 2 - 2)$  bits, plus a valid bit.

Thus, the complete cache size is

$$= 2^{10} \times \{4 \times 32 + (64 - 10 - 2 - 2) + 1\} = 2^{10} \times 179 \text{ bits}$$

= 22.4 KB for a 16 KB cache.

Consider a cache with 64 blocks and a block size of 16 bytes. To what block number does byte address 1200 map?

address of the block is =  $\text{floor} [(\text{Byte address}) / (\text{Bytes per block})] = \text{floor}(1200/16) = 75$

Block-id in the cache = (Block address) modulo (number of blocks in the cache)

$75 \bmod 64 = 11^{\text{th}}$  block

1200, 1201, ....1215

block address is the block containing all addresses between  $\text{floor} [ \text{Byte address} / \text{Bytes per block} ] * \text{Bytes per block} + (\text{no. of bytes in a block} - 1)$



# Example Problem

---

- *Consider a cache with 64 blocks and a block size of 16 bytes. What block number does byte address 1200 map to?*
- As block size = 16 bytes:  
byte address 1200  $\Rightarrow$  block address  $\lfloor 1200/16 \rfloor = 75$
- As cache size = 64 blocks:  
block address 75  $\Rightarrow$  cache block  $(75 \bmod 64) = 11$





# Cache Read Hit/Miss

---

- *Cache read hit:* no action needed
- *Instruction cache read miss:*
  1. *Send original PC value (current PC - 4, as PC has already been incremented in first step of instruction cycle) to memory*
  2. *Instruct main memory to perform read and wait for memory to complete access - stall on read*
  3. *After read completes write cache entry*
  4. *Restart instruction execution at first step to refetch instruction*
- *Data cache read miss:*
  - Similar to instruction cache miss
  - To reduce data miss penalty allow processor to execute instructions while waiting for the read to complete *until* the word is required - *stall on use* (why won't this work for instruction misses?)

# Cache Write Hit/Miss

## *Write-through* scheme

- on *write hit*: replace data in cache *and* memory with *every* write hit to avoid *inconsistency*
- on *write miss*: write the word into cache *and* memory – obviously no need to read missed word from memory!
- Write-through is slow because of always required memory write
  - performance is improved with a *write buffer* where words are stored while waiting to be written to memory – processor can continue execution until write buffer is full
  - when a word in the write buffer completes writing into main that buffer slot is freed and becomes available for future writes
  - DEC 3100 write buffer has 4 words
- *Write-back* scheme
  - write the data block *only* into the cache and *write-back* the block to main *only when* it is replaced in cache
  - more efficient than write-through, more complex to implement